Yili Yang, PhD
yyleli@gmail.com | +44 7490334233 | London, UK

12 August 2026

Hiring Team, Applied AI Research
Goldman Sachs

Dear Hiring Team,

I am writing to apply for the Applied AI Researcher position with the Applied AI Research team at Goldman Sachs.

For the past four years I have led a continental-scale deep-learning programme at a research institute, taking it from research design through to results that others depend on. A large part of that work has been the unglamorous half of applied research: getting model inference to run reliably and economically over very large datasets on Google Cloud Platform, rather than only in a notebook. Your posting's emphasis on advancing production machine-learning systems, and on maintaining production-ready code, describes the part of the job I have actually spent my time on.

Time-series work runs through the same period. I have built machine-learning models for time-series imputation where the observational record is incomplete, and for multi-domain prediction across heterogeneous signals. I hold a competitively awarded grant to generalise that into a single framework for multi-domain time-series prediction, which is the kind of reusable library-and-framework work your team describes rather than a one-off model.

I also care about the research being usable by other people. I designed and built a benchmark training dataset that is now a shared resource, built a zero-shot interactive segmentation workflow on top of the Segment Anything Model, and published a critical assessment of where foundation models fail rather than only where they succeed. Internally I lead an annual machine-learning workshop series, which is where most of my cross-team technical leadership has happened.

I should be direct about one requirement. I have not worked on quantitative investment workflows in equities, fixed income, or multi-asset strategies, and I have no financial-domain experience to offer. What I would bring instead is deep experience of machine learning under exactly the conditions your posting mentions as difficult: noisy and incomplete observational data, severe class imbalance, models that must run at scale and keep running, and results that get scrutinised hard by domain experts who did not build them. If the team needs someone who already knows the asset classes, I am not that person; if it needs a researcher who can learn a domain and be trusted with production systems in it, I would like to be considered.

I would welcome the chance to discuss the role and am happy to talk through any of the above in more detail. Thank you for your time and consideration.

Yours sincerely,

Yili Yang, PhD